

Run It Yourself.

Why local AI matters, what to expect when you run it,
and which machine to start with.

```
$ whoami → free■
```

The Labs want to slow down.

- ▶ **1,386 employees** sign “Pacing the Frontier” – coordinated pause
- ▶ **Amodei’s 3-step plan** – auditors first, then coordination
- ▶ **From** OpenAI · Anthropic · Google DeepMind · Meta
- ▶ **The pause costs them nothing** – it slows you, not them

Pacing the Frontier (Jul 2026, 1,386 signatories) · darioamodei.com “We Must Pace the Frontier” (Sep 2026)

Meanwhile, your plan got **smaller**.

~24 h

usable Claude Pro per month – paying
subscriber's log

2–2.5 d

until a \$200 Max plan hits its weekly cap

1.7–5×

faster quota burn after the Aug 17 reset

~5 h

a Codex plan's whole weekly allowance, gone

UNSTOPPABLE

weights don't answer to a pause

UNLIMITED

no begging on X for limit resets

ALIGNED

a model that answers to you, not to a ToS

```
$ cd ~/local && git log --oneline 2026
```

The **pace**

Six weeks of open-weight releases. Nobody asked permission.

2026, the year of local AI



Seven frontier-grade open releases in **six weeks**.

weights published → serving, on hardware you own

9 h 41 m

GLM-5.3-Flash NVFP4

16 h

DeepSeek-V4.1-Flash

same day

Qwen3.8-Flash-Next

\$ You can't pace what already shipped.■

Open weights ≠ open source.

Downloadable is not the same as free to use.

Qwen3.8-27B	Apache-2.0	OSI-approved – take it, ship it, sell it
GLM-5.3-Flash	MIT	OSI-approved – the Flash tier is the freer one
DeepSeek-V4.1-Flash	MIT	OSI-approved – same terms
GLM-5.3	custom “glm-5.3”	not OSI – MaaS over \$10B/yr needs Z.ai’s review
Kimi K3	custom “Kimi K3”	not OSI – MaaS over \$20M/yr needs a separate deal

```
$ lsblk --tiers
```

Pick your machine

Five tiers, two models. Measured in tokens/s – single-stream decode, 2026 community benchmarks.

MacBook Pro · M5 Pro, 48 GB

74 tok/s

- live – streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours – no usage email, no terms of service in the loop. every token local, every token yours.■

model ▶ Qwen 3.8–27B – Opus 4.6–class intelligence

Frontier-class coding and vision in a machine that fits a backpack. No CUDA, no daemon, no datacenter. Splash ships Metal kernels compiled for this model and a DFlash 2 draft trained for it – 282 ms to first token on a cached 32K context.

Splash 1.0 SPEED-Bench (incoai) on M5 Pro (16c), 48 GB · retrieved 2026-09-18

RTX 3090 - 24 GB

138 tok/s

- live – streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours – no usage email, no terms of service in the loop. every token local, every token yours.■

model ▶ Qwen 3.8–27B – Opus 4.6–class intelligence

A six-year-old used card, still the community's benchmark box. August's EXL3/W4A16 + DFlash2 kits rewrote the 24 GB recipe: 133–138 tok/s single-stream on reproduction-heavy prompts, 942 aggregate ×64 – no offload.

MiaAI-Lab EXL3 kit · 3090 stacks · retrieved 2026-09-17

GFTS 2026 | @sethforprivacy

11 / 15

2× DGX Spark · 256 GB

62.9 tok/s

- live – streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours – no usage email, no terms of service in the loop. every token local, every token yours.■

model ▶ GLM-5.3-Flash 320B – 0pus 4.8-class intelligence

12× the 3090's model, on two desk tiles instead of one card: 62.9 tok/s structured, 32 on prose, 146 aggregate at 4 streams, with a 1.75M-token KV pool.

github.com/MiaAI-Lab/GLM-5.3-Flash-EXL3-2x-DGX-Sparks (2026-08-28) · retrieved 2026-09-17

Mac Studio Ultra · 256 GB

51 tok/s

- live – streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours – no usage email, no terms of service in the loop. every token local, every token yours.■

model ▶ GLM-5.3-Flash 320B – 0pus 4.8-class intelligence

Sparse MoE does the heavy lifting – 321B of weights, 18B read per token. 25–32 tok/s plain, 51 with a drafter. Slow prefill is the tax.

2× RTX PRO 6000 - 192 GB

145.5 tok/s

- live – streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours – no usage email, no terms of service in the loop. every token local, every token yours.■

model ▶ GLM-5.3-Flash 320B – 0pus 4.8-class intelligence

Nothing offloaded. 6.2K tok/s cold prefill at 32–64K, 5/5 GSM8K drafter acceptance. The single-user endgame.

hf.co/brandonmusic/GLM-5.3-Flash-tr3-4bpw · retrieved 2026-09-17

GFTS 2026 | @sethforprivacy

14 / 15

Local AI is freedom tech.

```
$ splash serve --model incoai/Qwen3.8-27B-Splash
```

One 27B model, on the laptop you already own. Start this week.■