

GLOBAL FREEDOM TECH SUMMIT

Run It Yourself.

Why local AI matters, what to expect when you run it,
and which machine to start with.

```
$ whoami → free
```

```
local-ai --status ok
```

01 / 14

01 · THE ASK

The labs want to slow down.

- ▶ **1,386 employees** sign “Pacing the Frontier” – coordinated pause
- ▶ **Amodei’s 3-step plan** – norms, audits, global regulation
- ▶ **From** OpenAI · Anthropic · Google DeepMind · Meta
- ▶ **The pause costs them nothing** – it slows you, not them

`pacingthefrontier.com · darioamodei.com “We Must Pace the Frontier” · Sept 2026`

`local-ai --status ok`

02 / 14

01 · THE CATCH

Meanwhile, your plan got **smaller**.

~24 h

usable Claude Pro per month – paying
subscriber's log

2–2.5 d

until a \$200 Max plan hits its weekly cap

1.7–5×

faster quota burn after the Aug 17 reset

~5 h

a Codex plan's whole weekly allowance, gone

`anthropics/claude-code` issues #11810 · #65678 · #87419 · `community.openai.com` Codex threads

`local-ai --status` ok

03 / 14

UNSTOPPABLE weights don't answer to a pause

UNLIMITED no begging on X for limit resets

ALIGNED a model that answers to you, not to a ToS

```
$ cd ~/local && git log --oneline 2026
```

The **pace**

Six months of open-weight releases. Nobody asked permission.

2026, in six commits

JUL 31	AUG 14	AUG 14	AUG 26	SEP	SEP
V4-Flash 0731	Qwen 3.8- 27B	GLM-5.3	GLM-5.3- Flash	Qwen3.8- Flash-Next	DeepSeek- V4.1-Flash
official release · vLLM, llama.cpp, MLX day-one	dense · Apache- 2.0 · 262K ctx – 2.4T-A95B Max hours from GLM- 5.3	zai-org flagship – the open drop of the day	320B / 18B active · 1M ctx · MIT weights shipped	125B MoE · 1M ctx · vision – 37 tok/s on one Spark	flash-line refresh – day- one engine support

02 · WHY IT MATTERS

Six frontier-grade open releases in **six months.**

\$ You can't pace what already shipped.■

```
$ lsblk --tiers
```

Pick your machine

Five tiers, one model each. Measured in tokens/s – single-stream decode, 2026 community benchmarks.

```
local-ai --status ok
```

08 / 14

03 · TIER 01

MacBook Pro · M5 Max, 64 GB

61.6 tok/s

- live — streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours — no usage email, no terms of service in the loop. every token local, every token ■

model ▶ Qwen 3.8-27B — Opus 4.6-class intelligence

Frontier-class coding and vision in a machine that fits a backpack. No CUDA, no daemon, no datacenter. The engine, not the silicon, moved: 1.7× on the same laptop.

oMLX M5 Max (40c) 4-bit + MTP k=3: 61.6 · AWQ 5bpw: 59.4 · 5-bit no-MTP: 35.5 tok/s

local-ai --status ok

09 / 14

03 · TIER 02

RTX 3090 · 24 GB

138 tok/s

● live – streamed at this tier’s measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours – no usage email, no terms of service in the loop. every token local, every token yours.

model ▶ Qwen 3.8–27B – Opus 4.6–class intelligence

A five-year-old used card, still the community’s benchmark box. August’s EXL3/W4A16 + DFlash2 kits rewrote the 24 GB recipe: 133–138 tok/s chat, 942 aggregate ×64 – no offload.

github.com/MiaAI-Lab EXL3+DFlash2 kit · 3090 stacks – 133–138 tok/s single-stream, 942 at 64-way

local-ai --status ok

03 · TIER 03

2× DGX Spark · 256 GB

62.9 tok/s

- live – streamed at this tier’s measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours – no usage email, no terms of service in the loop. every token local, every token yours■

model ▶ GLM-5.3-Flash 320B – 0pus 4.8-class intelligence

12× the 3090’s model at the same token speed: 62.9 tok/s structured, 32 prose, 146 aggregate at 4 streams, on a 1M-token context – a team endpoint on two desk tiles.

github.com/MiaAI-Lab/GLM-5.3-Flash-EXL3-2x-DGX-Sparks – 62.9 tok/s ×1 · 146.5 agg ×4 · 1.75M-token KV pool (2026-08-28)

local-ai --status ok

11 / 14

03 · TIER 04

Mac Studio Ultra · 256 GB

51 tok/s

● live — streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is st

model ▶ GLM-5.3-Flash 320B — Opus 4.8-class intelligence

819 GB/s of bandwidth does the speculation heavy lifting: 25-32 tok/s plain, 51 with a drafter. Slow prefill is the tax.

llmcheck.net M3/M4 Ultra 25-32 · hf.co/grant-ai MLX-4bit DFlash2: 51.2 tok/s

local-ai --status ok

12 / 14

03 · TIER 05

2× RTX PRO 6000 · 192 GB

145.5 tok/s

- live — streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours — no usage email, no terms of service in the loop. every token local, every token yours.

model ▶ GLM-5.3-Flash 320B — 0pus 4.8-class intelligence

Nothing offloaded. 6.2K tok/s cold prefill, 262K context, 5/5 GSM8K drafter acceptance. The single-user endgame.

hf.co/brandonmusic/GLM-5.3-Flash-tr3-4bpw — 145.5 tok/s C1 decode, receipts in repo

local-ai --status ok

13 / 14

Local AI is freedom tech.

```
$ ollama run qwen3.8:27b-mlx
```

```
One 27B model, on the laptop you already own. Start this week.■
```