

GLOBAL FREEDOM TECH SUMMIT

Run It Yourself.

Why local AI matters, what to expect when you run it,
and which machine to start with.

```
$ whoami → free
```

```
local-ai --status ok
```

01 / 14

01 · THE ASK

The labs want to slow down.

- ▶ **1,386 employees** sign “Pacing the Frontier” – coordinated pause
- ▶ **Amodei’s 3-step plan** – norms, audits, global regulation
- ▶ **From** OpenAI · Anthropic · Google DeepMind · Meta
- ▶ **The pause costs them nothing** – it slows you, not them

pacingthefrontier.com · darioamodei.com “We Must Pace the Frontier” · Sept 2026

local-ai --status ok

02 / 14

01 · THE CATCH

Meanwhile, your plan got **smaller**.

~24 h

usable Claude Pro per month – paying
subscriber's log

2–2.5 d

until a \$200 Max plan hits its weekly cap

1.7–5×

faster quota burn after the Aug 17 reset

~5 h

a Codex plan's whole weekly allowance, gone

`anthropics/claude-code` issues #11810 · #65678 · #87419 · `community.openai.com` Codex threads

`local-ai --status` ok

03 / 14

UNSTOPPABLE weights don't answer to a pause

UNLIMITED no begging on X for limit resets

ALIGNED a model that answers to you, not to a ToS

```
$ cd ~/local && git log --oneline 2026
```

The **pace**

Six months of open-weight releases. Nobody asked permission.

2026, in six commits

JUL 31 V4-Flash 0731 official release · vLLM, llama.cpp, MLX day-one	AUG 14 Qwen 3.8-27B dense · Apache-2.0 · 262K ctx – 2.4T-A95B Max hours from GLM-5.3	AUG 14 GLM-5.3 zai-org flagship – the open drop of the day	AUG 26 GLM-5.3-Flash 320B / 18B active · 1M ctx · MIT weights shipped	SEP Qwen3.8-Flash-Next 125B MoE · 1M ctx · vision – 37 tok/s on one Spark	SEP DeepSeek-V4.1-Flash flash-line refresh – day-one engine support
---	---	---	--	--	--

02 · WHY IT MATTERS

Six frontier-grade open releases in **six months.**

\$ You can't pace what already shipped.

03 · TIER 02

RTX 3090 · 24 GB

138 tok/s

● live — streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours — no usage email, no terms of service in the loop. every token local, every token yours.

model ▶ Qwen 3.8-27B — Opus 4.6-class intelligence

A five-year-old used card, still the community's benchmark box. August's EXL3/W4A16 + DFlash2 kits rewrote the 24 GB recipe: 133-138 tok/s chat, 942 aggregate x64 — no offload.

github.com/MiaAI-Lab EXL3+DFlash2 kit · 3090 stacks — 133-138 tok/s single-stream, 942 at 64-way

local-ai --status ok

10 / 14

03 · TIER 03

2× DGX Spark · 256 GB

62.9 tok/s

- live — streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. pa

model ▶ GLM-5.3-Flash 320B — 0pus 4.8-class intelligence

12× the 3090's model at the same token speed: 62.9 tok/s structured, 32 prose, 146 aggregate at 4 streams, on a 1M-token context — a team endpoint on two desk tiles.

github.com/MiaAI-Lab/GLM-5.3-Flash-EXL3-2x-DGX-Sparks — 62.9 tok/s ×1 · 146.5 agg ×4 · 1.75M-token KV pool (2026-08-28)

local-ai --status ok

11 / 14

03 · TIER 04

Mac Studio Ultra · 256 GB

51 tok/s

- live — streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, █

model ▶ GLM-5.3-Flash 320B — Opus 4.8-class intelligence

819 GB/s of bandwidth does the speculation heavy lifting: 25-32 tok/s plain, 51 with a drafter. Slow prefill is the tax.

llmcheck.net M3/M4 Ultra 25-32 · hf.co/grant-ai MLX-4bit DFlash2: 51.2 tok/s

local-ai --status ok

12 / 14

03 · TIER 05

2× RTX PRO 6000 ▪ 192 GB

145.5 tok/s

● live — streamed at this tier's measured speed

no queue. no quota meter. no midnight reset. just weights, silence, and a model that answers to you. it reads the repo, drafts the patch, reviews its own diff, and asks for nothing between the first token and the last. past midnight it is still warm, still consistent, still yours — no usage email, no terms of service in the loop. every token local, every token yours.

model ▶ GLM-5.3-Flash 320B — 0pus 4.8-class intelligence

Nothing offloaded. 6.2K tok/s cold prefill, 262K context, 5/5 GSM8K drafter acceptance. The single-user endgame.

hf.co/brandonmusic/GLM-5.3-Flash-tr3-4bpw — 145.5 tok/s C1 decode, receipts in repo

local-ai --status ok

13 / 14

Local AI is freedom tech.

```
$ ollama run qwen3.8:27b-mlx
```

```
One 27B model, on the laptop you already own. Start this week.
```

Local AI is freedom tech.

```
$ ollama run qwen3.8:27b-mlx
```

```
One 27B model, on the laptop you already own. Start this week.█
```

Local AI is freedom tech.

```
$ ollama run qwen3.8:27b-mlx
```

```
One 27B model, on the laptop you already own. Start this week.
```